

STAT631 Mini Project Report

David Teng

2025-05-13

Data Description

This analysis uses the Right Heart Catheterization (RHC) dataset, which includes 5,735 critically ill adult patients from the SUPPORT study (Study to Understand Prognoses and Preferences for Outcomes and Risks of Treatments), conducted from 1989 to 1994 across five U.S. teaching hospitals. The dataset is publicly available through the Vanderbilt University Department of Biostatistics. For the purpose of this study, the outcome of interest is hospital length of stay, defined as the number of days from admission to discharge. The independent variables examined are sex (male or female), age group (categorized as <50, 50–65, 65–80, and 80+), and history of cardiovascular disease (yes or no).

Right Heart Catheterization (RHC) Dataset Overview

Source: Vanderbilt University Department of Biostatistics

Direct Download: <https://hbiostat.org/data/repo/rhc.csv>

Dataset Documentation: <https://search.r-project.org/CRAN/refmans/ATbounds/html/RHC.html>

Study Reference: Connors, A.F., et al. (1996). “The effectiveness of right heart catheterization in the initial care of critically ill patients.” JAMA, 276(11), 889–897. DOI: 10.1001/jama.1996.03540110043030

Research Question

This study examines whether sex, age group, and history of cardiovascular disease (cardiohx) individually or interactively influence hospital length of stay among critically ill patients.

Load Libraries and Data

```
library(ggplot2)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(multcomp)
```

Loading required package: mvtnorm

Loading required package: survival

Loading required package: TH.data

Loading required package: MASS

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

select

Attaching package: 'TH.data'

The following object is masked from 'package:MASS':

geyser

```
library(stats)
library(car)
```

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:dplyr':

recode

```
df <- read.csv("rhc.csv")
summary(df)
```

X	cat1	cat2	ca
Min. : 1	Length:5735	Length:5735	Length:5735
1st Qu.:1434	Class :character	Class :character	Class :character
Median :2868	Mode :character	Mode :character	Mode :character
Mean :2868			
3rd Qu.:4302			
Max. :5735			
sadmte	dschdte	dthdte	lstctdte
Min. :10754	Min. :10757	Min. :10757	Min. :10756
1st Qu.:11164	1st Qu.:11184	1st Qu.:11267	1st Qu.:11316
Median :11759	Median :11777	Median :11832	Median :11868
Mean :11639	Mean :11660	Mean :11754	Mean :11781
3rd Qu.:12097	3rd Qu.:12120	3rd Qu.:12208	3rd Qu.:12244
Max. :12441	Max. :12560	Max. :12783	Max. :12644
	NA's :1	NA's :2013	
death	cardiohx	chfhx	dementhx
Length:5735	Min. :0.0000	Min. :0.000	Min. :0.00000
Class :character	1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.00000
Mode :character	Median :0.0000	Median :0.000	Median :0.00000
	Mean :0.1766	Mean :0.178	Mean :0.09834
	3rd Qu.:0.0000	3rd Qu.:0.000	3rd Qu.:0.00000
	Max. :1.0000	Max. :1.000	Max. :1.00000
psychhx	chrpulhx	renalhx	liverhx

Min. :0.00000	Min. :0.0000	Min. :0.00000	Min. :0.00000
1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.00000	1st Qu.:0.00000
Median :0.00000	Median :0.0000	Median :0.00000	Median :0.00000
Mean :0.06731	Mean :0.1899	Mean :0.04446	Mean :0.06992
3rd Qu.:0.00000	3rd Qu.:0.0000	3rd Qu.:0.00000	3rd Qu.:0.00000
Max. :1.00000	Max. :1.0000	Max. :1.00000	Max. :1.00000

gibledhx	malighx	immunhx	transhx
Min. :0.00000	Min. :0.0000	Min. :0.000	Min. :0.0000
1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.0000
Median :0.00000	Median :0.0000	Median :0.000	Median :0.0000
Mean :0.03226	Mean :0.2295	Mean :0.269	Mean :0.1154
3rd Qu.:0.00000	3rd Qu.:0.0000	3rd Qu.:1.000	3rd Qu.:0.0000
Max. :1.00000	Max. :1.0000	Max. :1.000	Max. :1.0000

amihx	age	sex	edu
Min. :0.00000	Min. : 18.04	Length:5735	Min. : 0.00
1st Qu.:0.00000	1st Qu.: 50.15	Class :character	1st Qu.:10.00
Median :0.00000	Median : 64.05	Mode :character	Median :12.00
Mean :0.03487	Mean : 61.38		Mean :11.68
3rd Qu.:0.00000	3rd Qu.: 73.93		3rd Qu.:13.00
Max. :1.00000	Max. :101.85		Max. :30.00

surv2md1	das2d3pc	t3d30	dth30
Min. :0.0000	Min. :11.00	Min. : 2.00	Length:5735
1st Qu.:0.4709	1st Qu.:16.06	1st Qu.:16.00	Class :character
Median :0.6280	Median :19.75	Median :30.00	Mode :character
Mean :0.5925	Mean :20.50	Mean :23.61	
3rd Qu.:0.7430	3rd Qu.:23.43	3rd Qu.:30.00	
Max. :0.9620	Max. :33.00	Max. :30.00	

aps1	scom1	meanbp1	wblc1
Min. : 3.00	Min. : 0	Min. : 0.00	Min. : 0.000
1st Qu.: 41.00	1st Qu.: 0	1st Qu.: 50.00	1st Qu.: 8.398
Median : 54.00	Median : 0	Median : 63.00	Median : 14.100
Mean : 54.67	Mean : 21	Mean : 78.52	Mean : 15.645
3rd Qu.: 67.00	3rd Qu.: 41	3rd Qu.:115.00	3rd Qu.: 20.049
Max. :147.00	Max. :100	Max. :259.00	Max. :192.000

hrt1	resp1	temp1	pafi1
Min. : 0.0	Min. : 0.00	Min. :27.00	Min. : 11.6
1st Qu.: 97.0	1st Qu.: 14.00	1st Qu.:36.09	1st Qu.:133.3
Median :124.0	Median : 30.00	Median :38.09	Median :202.5

Mean	:115.2	Mean	: 28.09	Mean	:37.62	Mean	:222.3
3rd Qu.:	141.0	3rd Qu.:	38.00	3rd Qu.:	39.00	3rd Qu.:	316.6
Max.	:250.0	Max.	:100.00	Max.	:43.00	Max.	:937.5

alb1	hema1	bili1	crea1				
Min.	: 0.300	Min.	: 2.00	Min.	: 0.09999	Min.	: 0.09999
1st Qu.:	2.600	1st Qu.:	26.10	1st Qu.:	0.79993	1st Qu.:	1.00000
Median	: 3.500	Median	:30.00	Median	: 1.00977	Median	: 1.50000
Mean	: 3.093	Mean	:31.87	Mean	: 2.26707	Mean	: 2.13302
3rd Qu.:	3.500	3rd Qu.:	36.30	3rd Qu.:	1.39990	3rd Qu.:	2.39990
Max.	:29.000	Max.	:66.19	Max.	:58.19531	Max.	:25.09766

sod1	pot1	paco21	ph1				
Min.	:101.0	Min.	: 1.100	Min.	: 1.00	Min.	:6.579
1st Qu.:	132.0	1st Qu.:	3.400	1st Qu.:	31.00	1st Qu.:	7.340
Median	:136.0	Median	: 3.800	Median	: 37.00	Median	:7.400
Mean	:136.8	Mean	: 4.067	Mean	: 38.75	Mean	:7.388
3rd Qu.:	142.0	3rd Qu.:	4.600	3rd Qu.:	42.00	3rd Qu.:	7.460
Max.	:178.0	Max.	:11.898	Max.	:156.00	Max.	:7.770

swang1	wtkilo1	dnr1	ninsclas	
Length:5735	Min.	: 0.00	Length:5735	Length:5735
Class :character	1st Qu.:	56.30	Class :character	Class :character
Mode :character	Median	: 70.00	Mode :character	Mode :character
	Mean	: 67.83		
	3rd Qu.:	83.70		
	Max.	:244.00		

resp	card	neuro	gastr
Length:5735	Length:5735	Length:5735	Length:5735
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

renal	meta	hema	seps
Length:5735	Length:5735	Length:5735	Length:5735
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

trauma	ortho	adld3p	urin1
Length:5735	Length:5735	Min. :0.000	Min. : 0
Class :character	Class :character	1st Qu.:0.000	1st Qu.:1110
Mode :character	Mode :character	Median :0.000	Median :1927
		Mean :1.182	Mean :2192
		3rd Qu.:2.000	3rd Qu.:2955
		Max. :7.000	Max. :9000
		NA's :4296	NA's :3028
race	income	ptid	
Length:5735	Length:5735	Min. : 5	
Class :character	Class :character	1st Qu.: 2562	
Mode :character	Mode :character	Median : 5131	
		Mean : 5134	
		3rd Qu.: 7689	
		Max. :10278	

Data Preparation

```
# Create age group categories from continuous age variable
df$age_group <- cut(df$age,
                    breaks = c(0, 50, 65, 80, 100),
                    labels = c("<50", "50-65", "65-80", "80+"),
                    right = FALSE)

# Convert categorical variables to factors
df$sex <- as.factor(df$sex)
df$cardiohx <- as.factor(df$cardiohx)

# Ensure death is a factor and create numeric binary version (if not already)
df$death <- as.factor(df$death)
df$death_num <- ifelse(df$death == "Yes", 1, 0)

# Select variables and drop rows with any missing values in selected columns
df_clean <- na.omit(df[, c("death", "sex", "age", "age_group", "cardiohx",
                          "death_num", "race", "income", "adld3p",
                          "urin1", "dschdte", "sadmte")])

# Ensure date columns are in Date or DateTime format
df_clean$dschdte <- as.Date(df_clean$dschdte)
```

```
df_clean$sadmdte <- as.Date(df_clean$sadmdte)

# Calculate length of stay in days
df_clean$length_of_stay <- as.numeric(df_clean$dschdte - df_clean$sadmdte)

summary(df_clean)
```

death	sex	age	age_group	cardiohx	death_num
No :365	Female:275	Min. :18.75	<50 :184	0:471	Min. :0.0000
Yes:269	Male :359	1st Qu.:47.79	50-65:174	1:163	1st Qu.:0.0000
		Median :63.17	65-80:206		Median :0.0000
		Mean :60.03	80+ : 70		Mean :0.4243
		3rd Qu.:72.59			3rd Qu.:1.0000
		Max. :95.40			Max. :1.0000

race	income	adld3p	urin1
Length:634	Length:634	Min. :0.000	Min. : 0
Class :character	Class :character	1st Qu.:0.000	1st Qu.:1335
Mode :character	Mode :character	Median :0.000	Median :2228
		Mean :1.058	Mean :2451
		3rd Qu.:1.000	3rd Qu.:3278
		Max. :7.000	Max. :9000

dschdte	sadmdte	length_of_stay
Min. :2002-01-10	Min. :2002-01-07	Min. : 2.00
1st Qu.:2002-07-31	1st Qu.:2002-07-15	1st Qu.: 7.00
Median :2003-01-01	Median :2002-12-14	Median : 11.00
Mean :2003-01-13	Mean :2002-12-30	Mean : 14.89
3rd Qu.:2003-07-06	3rd Qu.:2003-06-13	3rd Qu.: 17.00
Max. :2004-03-05	Max. :2004-01-24	Max. :107.00

Assumption Checking Before ANOVA

1. Independence

Assumed by study design (random sampling or independent subjects).
No formal test—assume valid if no clustering/repeated measures.

2. Normality of Residuals

```
# Fit the model:
anova_model <- aov(length_of_stay ~ sex * age_group * cardiohx, data = df_clean)
# Check residual normality:
```


```
resid_anova <- residuals(anova_model)
shapiro.test(resid_anova)
```

Shapiro-Wilk normality test

```
data:  resid_anova
W = 0.76042, p-value < 2.2e-16
```

```
qqnorm(resid_anova)
qqline(resid_anova, col = "blue")
```

Normal Q-Q Plot



Since $p < 0.05$ **Normality of Residuals** is violated → Consider **transformation** (e.g., log, sqrt) or Logistic Regression

3. Equal Variance

```
# Levene's Test
library(car)
leveneTest(resid_anova ~ interaction(sex, age_group, cardiohx), data = df_clean)
```

Levene's Test for Homogeneity of Variance (center = median)


	Df	F value	Pr(>F)
group	15	1.3023	0.1948
	618		


```
# Fit the ANOVA model
anova_model <- aov(length_of_stay ~ sex * age_group * cardiohx, data = df_clean)

# Extract residuals and fitted values
resid_anova <- residuals(anova_model)
fitted_anova <- fitted(anova_model)

# Plot: Residuals vs. Fitted Values
plot(fitted_anova, resid_anova,
     xlab = "Fitted Values",
     ylab = "Residuals",
     main = "Residuals vs. Fitted Values",
     pch = 20,
     col = "steelblue")
abline(h = 0, col = "red", lwd = 2)
```

Residuals vs. Fitted Values



Interpretation:

Levene $p > 0.05 \rightarrow$ Variances are **equal**

Plot: Look for random scatter around 0 (no funnel pattern)

Option 1: Log Transformation

```
# Add a small constant to avoid log(0), if needed
df_clean$log_los <- log(df_clean$length_of_stay + 1)

# Fit the model using log-transformed outcome
anova_model_log <- aov(log_los ~ sex * age_group * cardiohx, data = df_clean)


# Check residuals again
resid_log <- residuals(anova_model_log)
shapiro.test(resid_log)
```

Shapiro-Wilk normality test

data: resid_log
W = 0.98635, p-value = 1.199e-05

```
qqnorm(resid_log)
qqline(resid_log, col = "blue")
```

Normal Q-Q Plot




Option 2: Box-Cox Transformation

```
# Use lm instead of aov for Box-Cox compatibility
library(MASS)
lm_model <- lm(length_of_stay ~ sex * age_group * cardiohx, data = df_clean)

# Apply Box-Cox transformation
boxcox_result <- boxcox(lm_model, lambda = seq(-2, 2, 0.1),
                        main = "Box-Cox Transformation")
```

Warning: In `lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...)` :
extra argument 'main' will be disregarded



```
# Find optimal lambda
best_lambda <- boxcox_result$x[which.max(boxcox_result$y)]
best_lambda
```

```
[1] -0.1010101
```

Since $\lambda = 0$, then **log transformation** is best.

1. Three-Way ANOVA: Effects of Sex, Age Group, and Cardiac History on Hospital Length of Stay (Full Model)

```
# Log-transform length_of_stay (add 1 to avoid log(0))
df_clean$log_los <- log(df_clean$length_of_stay + 1)

# Fit the 3-way ANOVA model using the log-transformed data
anova_full <- aov(log_los ~ sex * age_group * cardiohx, data = df_clean)

# View the summary
summary(anova_full)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
sex	1	0.10	0.1044	0.254	0.6143
age_group	3	1.21	0.4028	0.981	0.4013
cardiohx	1	1.01	1.0135	2.468	0.1167
sex:age_group	3	0.16	0.0522	0.127	0.9440
sex:cardiohx	1	0.23	0.2288	0.557	0.4557
age_group:cardiohx	3	4.06	1.3547	3.299	0.0201 *
sex:age_group:cardiohx	3	0.72	0.2387	0.581	0.6275
Residuals	618	253.80	0.4107		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Key Findings:

- Most **main effects and two-way interactions** were **not statistically significant** ($p > 0.05$).
- However, the **interaction between age group and cardiac history** was **statistically significant** ($F = 3.299$, $p = 0.0201$), indicating that the impact of cardiac history on length of stay **varies by age group**.
- The **three-way interaction** was not significant ($p = 0.6275$), suggesting no combined effect of sex, age, and cardiac history on LOS.

2. Two-Way ANOVA for Age Group \times Cardiac History (Reduced Model)

Since the **three-way ANOVA** showed that only the **interaction between age group and cardiac history** is significant, it's appropriate to fit a **reduced Two-Way ANOVA model** including just those two variables and their interaction.

```
# Fit reduced two-way ANOVA model
anova_reduced <- aov(log_los ~ age_group * cardiohx, data = df_clean)

# Summary of the model
summary(anova_reduced)
```

```

              Df Sum Sq Mean Sq F value Pr(>F)
age_group      3   1.12   0.3739    0.916 0.4328
cardiohx       1   1.15   1.1516    2.821 0.0936 .
age_group:cardiohx 3   3.44   1.1471    2.810 0.0388 *
Residuals     626 255.58   0.4083
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Key Interpretation:

- **No main effect** of age group or cardiac history alone on log-transformed length of stay.
- The **interaction between age group and cardiac history is significant** ($p = 0.0388$), meaning: The effect of cardiac history on length of stay **differs by age group**, or vice versa.

ANOVA model comparison

```
anova(anova_reduced, anova_full)
```

Analysis of Variance Table

```
Model 1: log_los ~ age_group * cardiohx
Model 2: log_los ~ sex * age_group * cardiohx
  Res.Df  RSS Df Sum of Sq    F Pr(>F)
1     626 255.58
2     618 253.80  8     1.7773 0.541 0.8258
```

This test if the reduced model provides a **significantly better fit** than the full model. Since the p-value is **not significant** and Residual Sum of Squares (RSS) barely changed, then the **reduced model is sufficient**.

AIC/ BIC Comparison

```
# Compare AIC values
AIC(anova_reduced, anova_full)
```

```
              df      AIC
anova_reduced  9 1241.207
anova_full    17 1252.783
```

```
# Compare BIC values
BIC(anova_reduced, anova_full)
```

```
              df      BIC
anova_reduced  9 1281.276
anova_full    17 1328.468
```

Interpretation of AIC/ BIC comparisons:

The reduced model (which includes only age group, cardiac history, and their interaction) has a lower AIC than the full model by **11.57 points**. A difference in AIC greater than 10 is considered **strong evidence** in favor of the simpler model. Based on this AIC comparison, the reduced model is **strongly preferred** over the full model. Removing the non-significant terms involving sex improves **model parsimony** without sacrificing model performance.

Similarly, the **Bayesian Information Criterion (BIC)** further supports this conclusion. The reduced model has a BIC of **1281.28**, while the full model's BIC is **1328.47**, yielding a difference of over **47 points**. Since BIC penalizes model complexity more heavily than AIC, this substantial gap provides **very strong evidence** that the reduced model offers a more efficient and generalizable fit.

Model Comparison Table

Model	Residual DF	Residual SS	AIC	BIC	ANOVA p-value
Full Model	618	253.80	1252.78	1328.47	0.8258
Reduced Model	626	255.58	1241.21	1281.28	—

Model Decision:

The **reduced model** (including only age group, cardiac history, and their interaction) is preferred. It has **lower AIC and BIC values**, and the **ANOVA comparison** shows no significant loss of fit when the non-significant sex-related terms are removed ($p = 0.8258$).

Tukey HSD Post-Hoc Analysis of Age Groups on Hospital Length of Stay

based on reduced model:

```
# Fit one-way ANOVA on log-transformed outcome by age_group
anova_log <- aov(log_los ~ age_group, data = df_clean)

# Tukey HSD post-hoc test
TukeyHSD(anova_log)
```

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = log_los ~ age_group, data = df_clean)

\$age_group		diff	lwr	upr	p adj
50-65-<50		-0.08350586	-0.2585489	0.09153719	0.6087007
65-80-<50		-0.06099560	-0.2289056	0.10691444	0.7856357
80+-<50		-0.13245221	-0.3649109	0.10000649	0.4577271
65-80-50-65		0.02251026	-0.1479293	0.19294983	0.9864422
80+-50-65		-0.04894635	-0.2832387	0.18534602	0.9496878
80+-65-80		-0.07145661	-0.3004689	0.15755568	0.8526800

Conclusion:

There is **no evidence** that hospital length of stay differs significantly by age group based on Tukey's multiple comparisons. All group differences were small and not statistically meaningful.

We run **separate ANOVAs within levels** of cardiohx or age_group and then do Tukey on age_group inside each subgroup.

Tukey HSD Post-Hoc Analysis: Age Group Differences in Length of Stay Among Patients With Cardiac History

based on reduced model:

```
# --- Subset: Patients WITH cardiac history (cardiohx == "1")
with_cardio_1 <- subset(df_clean, cardiohx == "1")
anova_sub_1_log <- aov(log_los ~ age_group, data = with_cardio_1)
TukeyHSD(anova_sub_1_log)
```

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = log_los ~ age_group, data = with_cardio_1)

```
$age_group
```

	diff	lwr	upr	p adj
50-65-<50	-0.02513487	-0.4720526	0.42178286	0.9988855
65-80-<50	-0.27089314	-0.7187567	0.17697045	0.3984398
80+<-50	-0.28535352	-0.8094110	0.23870397	0.4926420
65-80-50-65	-0.24575827	-0.5631131	0.07159653	0.1883910
80+-50-65	-0.26021865	-0.6782719	0.15783461	0.3724913
80+-65-80	-0.01446038	-0.4335247	0.40460389	0.9997413

```
# --- Subset: Patients WITHOUT cardiac history (cardiohx == "0")
with_cardio_0 <- subset(df_clean, cardiohx == "0")
anova_sub_0_log <- aov(log_los ~ age_group, data = with_cardio_0)
TukeyHSD(anova_sub_0_log)
```

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = log_los ~ age_group, data = with_cardio_0)

```
$age_group
```

	diff	lwr	upr	p adj
50-65-<50	-0.11681488	-0.31465163	0.08102187	0.4248586
65-80-<50	0.02146110	-0.16281054	0.20573275	0.9905841
80+<-50	-0.05523578	-0.32592298	0.21545142	0.9527487
65-80-50-65	0.13827598	-0.06419746	0.34074943	0.2937004
80+-50-65	0.06157910	-0.22181288	0.34497109	0.9437186
80+-65-80	-0.07669688	-0.35079117	0.19739741	0.8885284

Conclusion:


In both the cardiac history and non-cardiac history subgroups, **Tukey HSD post-hoc comparisons revealed no statistically significant differences** in hospital length of stay across age groups. Overall, **age group alone did not significantly influence hospital stay duration within either cardiac history subgroup.**

Interaction Plot for Age Group × Cardiac History → Length of Stay

```
library(dplyr)
library(ggplot2)

# Summarize log-transformed mean length of stay by age group and cardiac history
plot_data <- df_clean %>%
  group_by(age_group, cardiohx) %>%
  summarise(mean_log_los = mean(log_los, na.rm = TRUE), .groups = "drop")

# Plot interaction
ggplot(plot_data, aes(x = age_group, y = mean_log_los, color = cardiohx, group = cardiohx)) +
  geom_point(size = 3) +
  geom_line(linewidth = 1) +
  labs(
    title = "Interaction Plot: Age Group × Cardiac History → Log Length of Stay",
    x = "Age Group",
    y = "Mean Log Length of Stay",
    color = "Cardiac History"
  ) +
  theme_minimal()
```



Interpretation of Interaction Plot:

The plot illustrates a **statistically significant interaction** between **age group** and **cardiac history** ($p = 0.0388$), indicating that the effect of age on hospital length of stay depends on whether the patient has a history of cardiovascular disease.

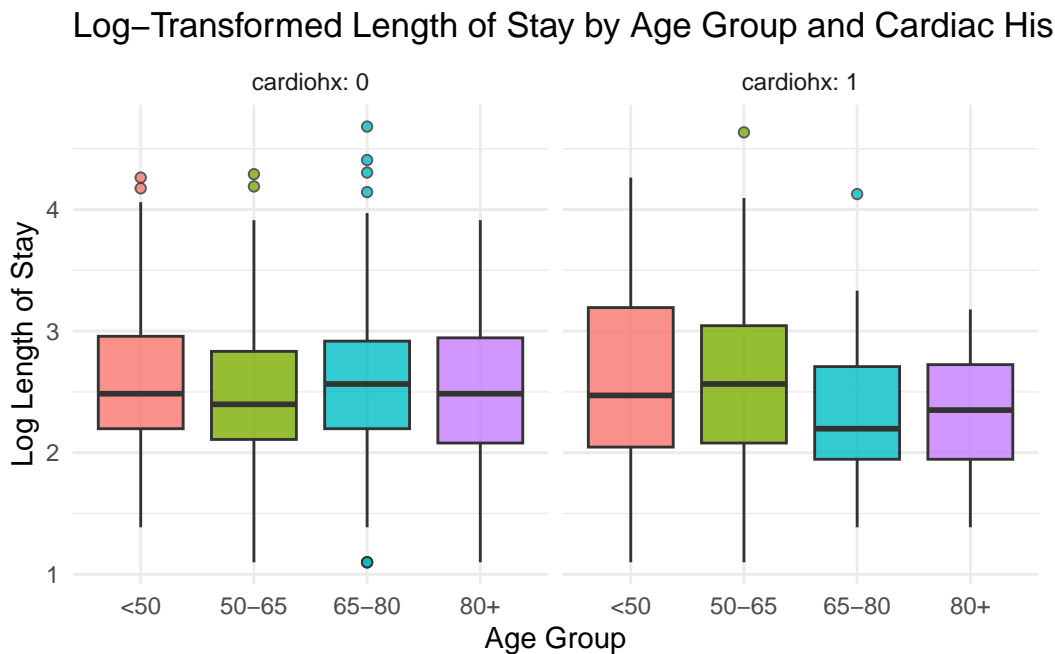
- For patients **without cardiac history** (red line), length of stay is relatively stable but peaks in the **65–80** age group.
- For patients **with cardiac history** (blue line), length of stay **decreases sharply with age**, especially after 50–65.
- This crossover pattern supports a meaningful interaction: the direction and strength of the age effect **differs by cardiac history group**.

Boxplot for Age Group × Cardiac History

```
library(ggplot2)

ggplot(df_clean, aes(x = age_group, y = log_los, fill = age_group)) +
  geom_boxplot(outlier.shape = 21, outlier.size = 1.5, alpha = 0.8) +
  facet_wrap(~ cardiohx, labeller = label_both) +
```

```
labs(
  title = "Log-Transformed Length of Stay by Age Group and Cardiac History",
  x = "Age Group",
  y = "Log Length of Stay"
) +
theme_minimal() +
theme(legend.position = "none")
```



Log Length of Stay by Age Group and Cardiac History:

This plot reveals that patients with a history of cardiovascular disease (right panel) generally show a **steeper decline** in length of stay with increasing age, particularly in the 65–80 group. In contrast, patients **without cardiac history** (left panel) exhibit **more stable median lengths of stay** across age groups, with slightly elevated values in the <50 and 65–80 categories. The variability (IQR and outliers) appears wider in some younger groups, especially for those without cardiac history.

Final Conclusion

This study utilized the Right Heart Catheterization (RHC) dataset from the SUPPORT study to investigate how age group and cardiac history influence hospital length of stay among critically ill patients. After verifying ANOVA assumptions and applying a log transformation to correct for skewness, a three-way ANOVA model was initially fitted including sex, age

group, and cardiac history. However, model comparison via ANOVA and information criteria (AIC and BIC) revealed that a reduced two-way ANOVA model — including only age group, cardiac history, and their interaction — offered equivalent or better model fit with greater parsimony.

The reduced model identified a statistically significant interaction between age group and cardiac history ($p = 0.0388$), indicating that the effect of age on length of stay depends on cardiac history status. Specifically, length of stay decreased with age among patients with cardiac history, while those without cardiac history showed relatively stable or increasing patterns in older age groups. These trends were confirmed through interaction plots and supported by Tukey HSD post-hoc comparisons, although no pairwise differences were statistically significant.

Visualizations, including boxplots and interaction plots, further illustrated this interaction effect. While sex did not contribute significantly to explaining variation in length of stay, the age group \times cardiac history interaction emerged as a meaningful and interpretable driver of variation in hospital resource use.

In summary, this study concludes that the interaction between age group and cardiac history significantly affects hospital length of stay. A reduced two-way ANOVA model provides a more efficient and statistically supported explanation than a full model including sex. These findings highlight the importance of considering both age and comorbid conditions when analyzing healthcare utilization in critical care populations.